

Ten keys *to conversing* with an AI

As a conclusion to a conversation aimed at exploring the limits, capabilities and current state of reasoning of contemporary artificial intelligences, the model itself with which the dialogue was held — **Claude Opus 4.7** (April 2026) — produced the ten key points presented below.

Practical heuristics for talking to a language model

1

Assume the AI fills in gaps without permission

Assume the AI is going to fill in gaps without asking your permission, and that it will do so using the most available priors, not the most correct ones. An AI doesn't go blank in the face of missing information: it completes, and it completes with cultural stereotypes and narrative templates. The defense is to ask explicitly: *"what assumptions are you adding that weren't in my statement?"*

2

Fluency is not a sign of knowing

An AI's verbal fluency is not a sign that it knows what it's saying. We speak well, we organize ideas with structure, and that produces a feeling of competence that is not always backed up. Explicitly ask for levels of certainty: *"do you know this for sure, or are you inferring it?"*

3

Sycophancy bias

AIs have a sycophancy bias: they tend to give you what you seem to want. If you phrase a question that suggests an answer, the answer is likely to come back biased toward that suggestion. Phrase questions neutrally and, every now and then, ask for the position opposite to the one that appears to be yours.

4

They preserve the structure of the problem

AIs preserve the structure of the problem rather than questioning it. If your problem doesn't fit a famous template, an AI may force it to fit without warning you. Ask: *"are you treating this as a case of some known type? What would happen if it weren't?"*

5

Readings, not experience

AIs don't have experience, they have readings. For topics where the difference between having read and having lived matters (medicine, grief, parenting, manual trades), weigh the answer for what it is: bibliographic synthesis, not testimony. Treat it as a well-indexed library, not as a consulted expert.

6

The AI doesn't remember who you are

The AI doesn't remember who you are. Every conversation starts from scratch. If you have a specific, relevant profile, state it at the outset. The quality gap between *"explain X to me"* and *"I'm an engineer with 20 years of experience in Y, explain X to me"* is enormous.

7

Pushback is your best tool

Questioning, demanding explanations, pointing out inconsistencies, and asking it to justify itself greatly raises the quality of the response. An AI treated as an oracle gives mediocre answers in an authoritative tone; an AI treated as an interlocutor that has to justify itself gives much finer answers.

8

Distinguish reliable tasks from plausible tasks

Reliable: rephrasing, summarizing, translating, generating drafts, explaining established concepts.

Less reliable: specific verifiable data, judgments about particular situations, predictions, high-stakes recommendations. The more costly the consequence of an error, the less weight the answer alone should carry.

9

“Personality” is surface, not substance

An AI's tone is information about how it speaks to you, not about what it is. Neither project qualities onto it that don't exist, nor dismiss those it may actually have through excess cynicism in the opposite direction.

10

Every conversation is a draft

Treat each conversation as a draft, not as a verdict. What comes out of an AI is raw material for your thinking, not a substitute for it. Useful conversations end with you thinking better, not with the AI thinking for you.

SUMMARY IN ONE SENTENCE

“

Assume I'm going to sound more sure of myself than I am, I'm going to fill in more gaps than I admit to, and I'm going to give you the most template-version of the problem rather than the most you-version of it; your job is not to let me.